

MC855 - Projeto em Sistemas de Computação

Apache Spark

Islene Calciolari Garcia

Instituto de Computação - Unicamp

Primeiro Semestre de 2016

Sumário

Revisão de MapReduce

ApacheSpark

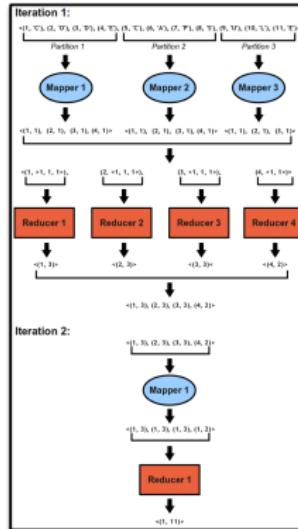
Transformações e Ações

Projeto 2



MapReduce: A Programming Model

- **MapReduce:**
Simplified Data Processing on Large Clusters
(published 2004)
- **Parallel and Distributed Algorithm:**
- Data Locality
- Fault Tolerance
- Linear Scalability



Fonte: Carol McDonald: An Overview of Apache Spark



MapReduce Basics

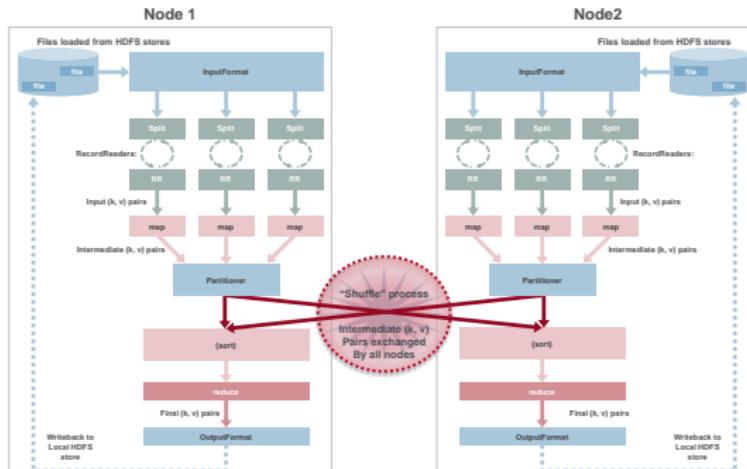
- Foundational model is based on a distributed file system
 - Scalability and fault-tolerance
- Map
 - Loading of the data and defining a set of keys
 - Many use cases do not utilize a reduce task
- Reduce
 - Collects the organized key-based data to process and output
- Performance can be tweaked based on known details of your source files and cluster shape (size, total number)



Fonte: Carol McDonald: An Overview of Apache Spark



MapReduce Execution and Data Flow

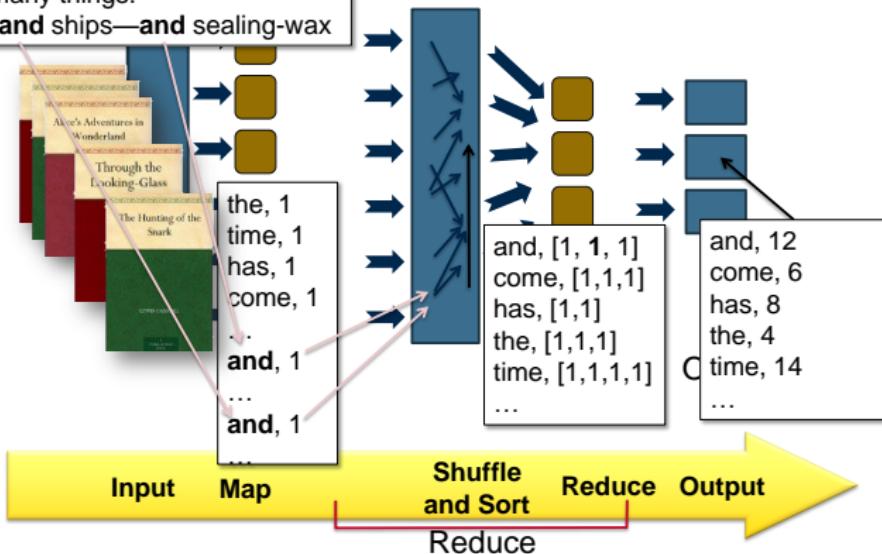


Fonte: Carol McDonald: An Overview of Apache Spark



MapReduce Example: Word Count

"The time has come," the Walrus said,
"To talk of many things:
Of shoes—and ships—and sealing-wax

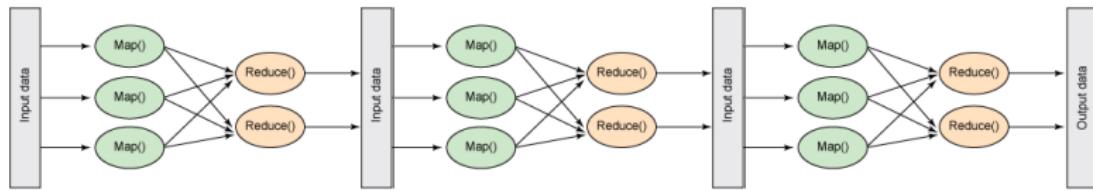


Fonte: Carol McDonald: An Overview of Apache Spark



MapReduce Processing Model

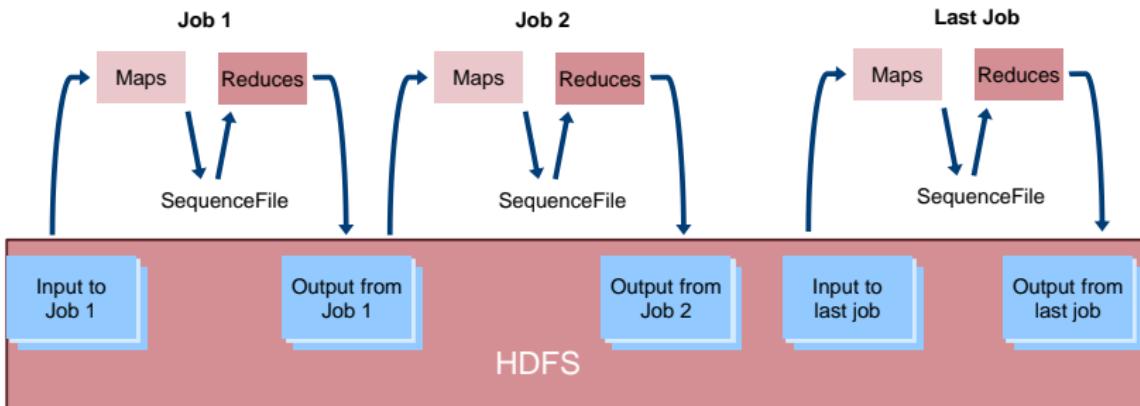
- Define mappers
- Shuffling is automatic
- Define reducers
- For complex work, chain jobs together
 - Use a higher level language or DSL that does this for you



Fonte: Carol McDonald: An Overview of Apache Spark



Typical MapReduce Workflows



Fonte: Carol McDonald: An Overview of Apache Spark



MapReduce: The Good

- Built in fault tolerance
- Optimized IO path
- Scalable
- Developer focuses on Map/Reduce, not infrastructure
- simple? API



Fonte: Carol McDonald: An Overview of Apache Spark



MapReduce: The Bad

- Optimized for disk IO
 - Doesn't leverage memory well
 - **Iterative** algorithms go through disk IO path again and again
- Primitive API
 - simple abstraction
 - Key/Value in/out
 - basic things like join
 - require extensive code
- Result often many files that need to be combined appropriately



Fonte: Carol McDonald: An Overview of Apache Spark

Apache Spark



©

2014 MapR Technologies



Fonte: Carol McDonald: An Overview of Apache Spark



Apache Spark

- Originally developed in 2009 in UC Berkeley's AMP Lab
- Fully open sourced in 2010 – now a Top Level Project at the Apache Software Foundation

spark.apache.org
github.com/apache/spark
user@spark.apache.org

The screenshot shows the official Apache Spark website. At the top is the Apache logo and the text "Apache Spark - Lightning-fast cluster computing". Below the header is a navigation bar with links for "Download", "Related Projects", "Documentation", "Community", and "FAQ". A main banner states "Apache Spark is a fast and general engine for large-scale data processing." To the left, a section titled "Speed" compares Hadoop and Spark's performance on logistic regression. A bar chart shows running times in seconds: Hadoop takes 110 seconds while Spark takes 0.9 seconds. Below this is a section titled "Ease of Use" with a code snippet demonstrating how to read a file in Java or Scala.

System	Running time (s)
Hadoop	110
Spark	0.9

```
file = spark.textFile("hdfs://...")  
file.flatMapLine => line.split(" ")  
.mapWord => (word, 1)  
.reduceByKey
```

Fonte: Carol McDonald: An Overview of Apache Spark



- Rich APIs in Java, Scala, Python
- Interactive shell
- **Fast to Run**
 - General execution graphs
 - In-memory storage

2-5x less code



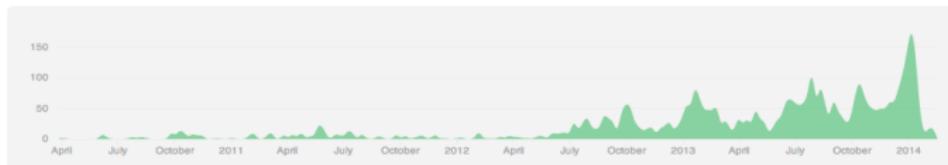
Fonte: Carol McDonald: An Overview of Apache Spark



The Spark Community

March 27th 2010 – February 15th 2014
Commits to master, excluding merge commits

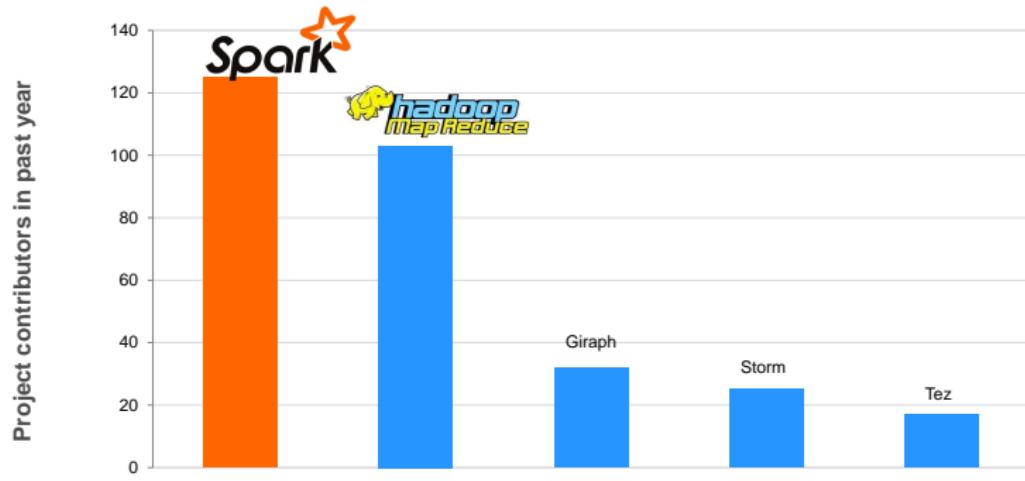
Contribution Type: **Commits** ▾



Fonte: Carol McDonald: An Overview of Apache Spark



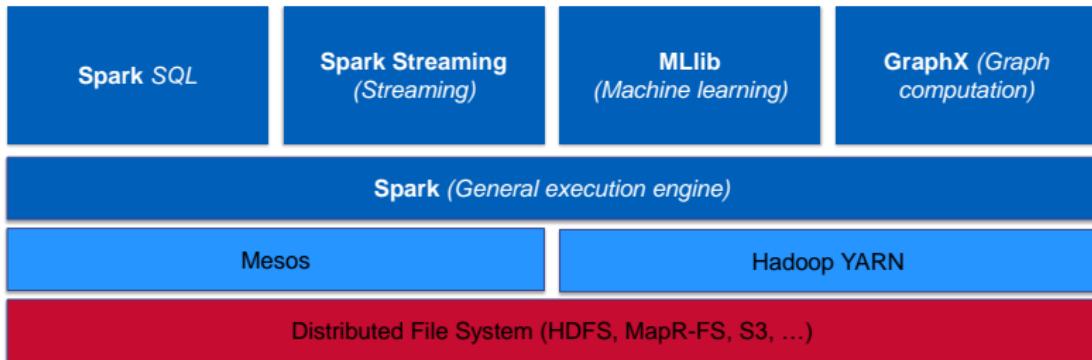
Spark is the Most Active Open Source Project in Big Data



Fonte: Carol McDonald: An Overview of Apache Spark



Unified Platform



Fonte: Carol McDonald: An Overview of Apache Spark



- Iterative Algorithms on large amounts of data
- Anomaly detection
- Classification
- Predictions
- Recommendations





Why Iterative Algorithms

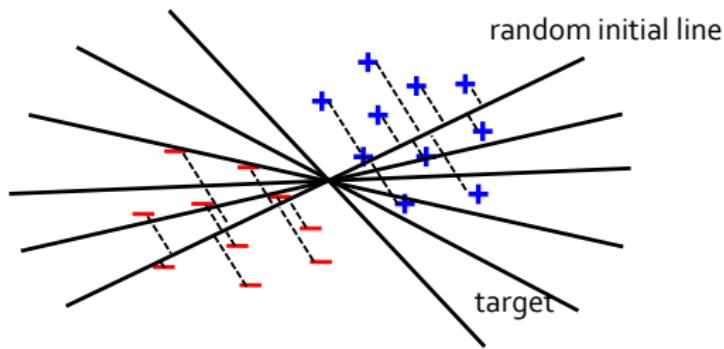
- Algorithms that need iterations
 - Clustering (K-Means, Canopy, ...)
 - Gradient descent (e.g., Logistic Regression, Matrix Factorization)
 - Graph Algorithms (e.g., PageRank, Line-Rank, components, paths, reachability, centrality,)
 - Alternating Least Squares ALS
 - Graph communities / dense sub-components
 - Inference (believe propagation)
 - ...





Example: Logistic Regression

- Goal: find best line separating two sets of points



Fonte: Carol McDonald: An Overview of Apache Spark



Logistic Regression

```
data = spark.textFile(...).map(readPoint).cache()

w = numpy.random.rand(D)

for i in range(iterations):
    gradient = data
        .map(lambda p: (1 / (1 + exp(-p.y * w.dot(p.x))))
                     * p.y * p.x)
        .reduce(lambda x, y: x + y)
    w -= gradient

print "Final w: %s" % w
```

Iteration!



Fonte: Carol McDonald: An Overview of Apache Spark



Data Sources

- Local Files
 - file:///opt/httpd/logs/access_log
- S3
- Hadoop Distributed Filesystem
 - Regular files, sequence files, any other Hadoop InputFormat
- HBase
- other NoSQL data stores



Fonte: Carol McDonald: An Overview of Apache Spark

How Spark Works

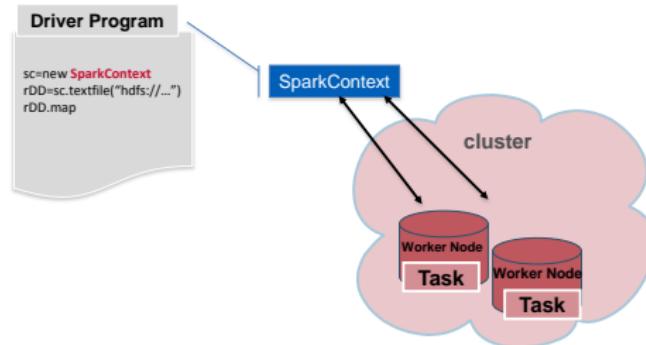


© 2014 MapR Technologies 

Fonte: Carol McDonald: An Overview of Apache Spark



Spark Programming Model



Fonte: Carol McDonald: An Overview of Apache Spark

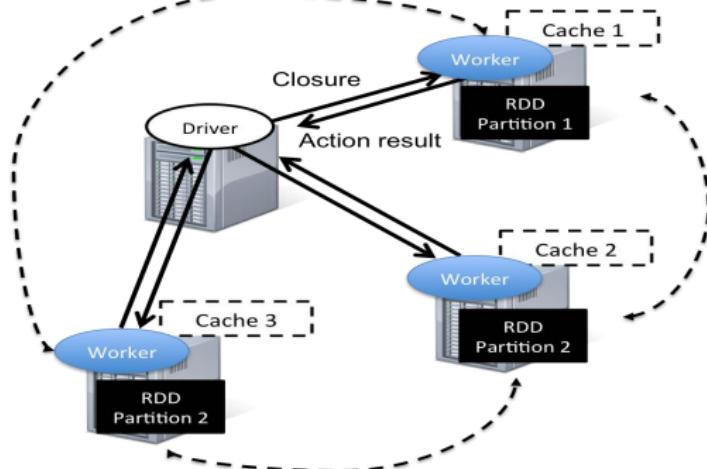


Resilient Distributed Datasets (RDD)

Spark revolves around RDDs

- Fault-tolerant
- read only collection of elements
- operated on in parallel
- Cached in memory
- Or on disk

http://www.cs.berkeley.edu/~matei/papers/2012/nsdi_spark.pdf



Fonte: Carol McDonald: An Overview of Apache Spark



Working With RDDs

```
textFile = sc.textFile("SomeFile.txt")
```

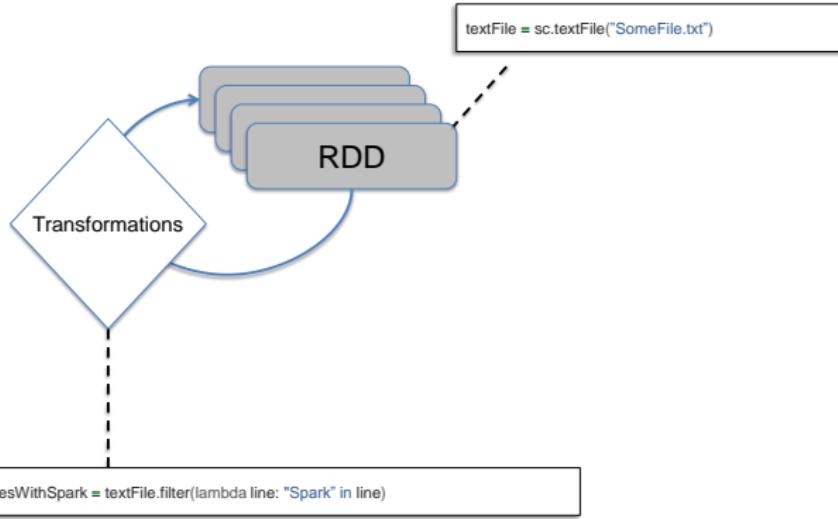
RDD



Fonte: Carol McDonald: An Overview of Apache Spark



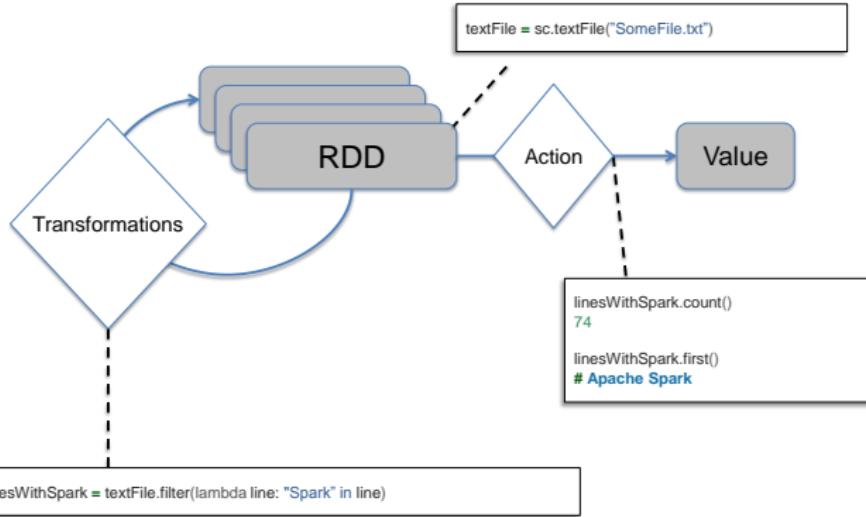
Working With RDDs



Fonte: Carol McDonald: An Overview of Apache Spark



Working With RDDs



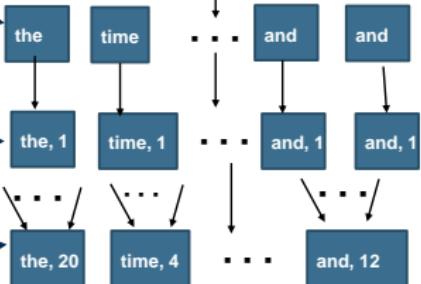
Fonte: Carol McDonald: An Overview of Apache Spark



Example Spark Word Count in Java

```
JavaRDD<String> input = sc.textFile(inputFile);
// Split each line into words
JavaRDD<String> words = input.flatMap(
    new FlatMapFunction<String, String>() {
        public Iterable<String> call(String x) {
            return Arrays.asList(x.split(" "));
        }
});
// Turn the words into (word, 1) pairs
JavaPairRDD<String, Integer> word1s == words.mapToPair(
    new PairFunction<String, String, Integer>(){
        public Tuple2<String, Integer> call(String x){
            return new Tuple2(x, 1);
        }
});
// reduce add the pairs by key to produce counts
JavaPairRDD<String, Integer> counts = word1s.reduceByKey(
    new Function2<Integer, Integer, Integer>(){
        public Integer call(Integer x, Integer y){
            return x + y;
        }
});
```

"The time has come," the Walrus said,
"To talk of many things:
Of shoes—and ships—and sealing-wax



Fonte: Carol McDonald: An Overview of Apache Spark



Example Spark Word Count in Scala

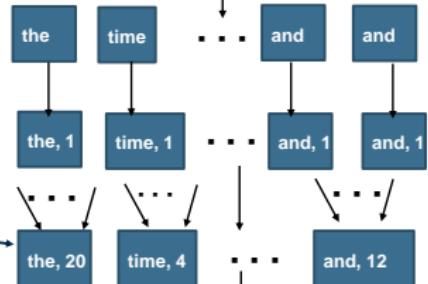
```
// Load our input data.  
val input = sc.textFile(inputFile)
```

```
// Split it up into words.  
val words = input.flatMap(line => line.split(" "))
```

```
// Transform into pairs and count.  
val counts = words  
  .map(word => (word, 1))  
  .reduceByKey{case (x, y) => x + y}
```

```
// Save the word count back out to a text file,  
counts.saveAsTextFile(outputFile)
```

"The time has come," the Walrus said,
"To talk of many things:
Of shoes—and ships—and sealing-wax



the, 20 time, 4 and, 12



Fonte: Carol McDonald: An Overview of Apache Spark

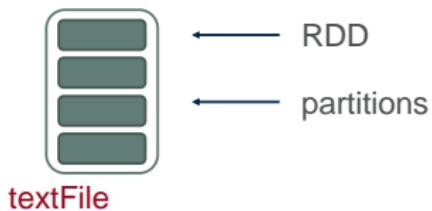


Example Spark Word Count in Scala

```
// Load input data.  
val input = sc.textFile(inputFile)
```

HadoopRDD

MapPartitionsRDD



64

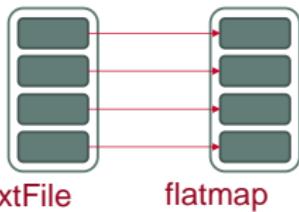
Fonte: Carol McDonald: An Overview of Apache Spark



Example Spark Word Count in Scala

```
// Load our input data.  
val input = sc.textFile(inputFile)  
// Split it up into words.  
val words = input.flatMap(line => line.split(" "))
```

HadoopRDD
MapPartitionsRDD
MapPartitionsRDD



65

Fonte: Carol McDonald: An Overview of Apache Spark



FlatMap



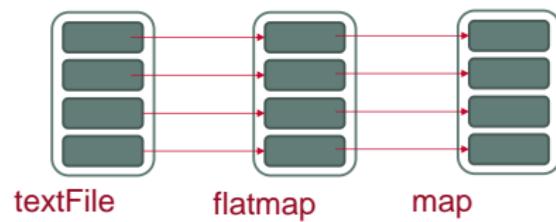
Fonte: Carol McDonald: An Overview of Apache Spark



Example Spark Word Count in Scala

```
val input = sc.textFile(inputFile)
val words = input.flatMap(line => line.split(" "))
// Transform into pairs
val counts = words.map(word => (word, 1))
```

HadoopRDD
MapPartitionsRDD
MapPartitionsRDD
MapPartitionsRDD



67



Fonte: Carol McDonald: An Overview of Apache Spark



Map

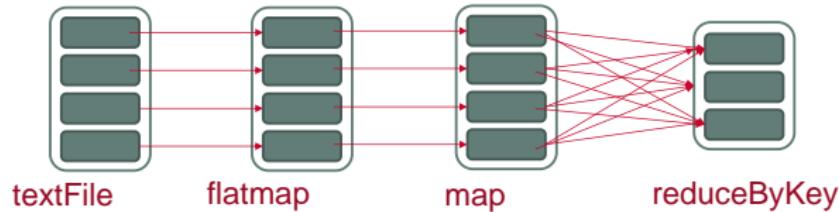


Fonte: Carol McDonald: An Overview of Apache Spark



Example Spark Word Count in Scala

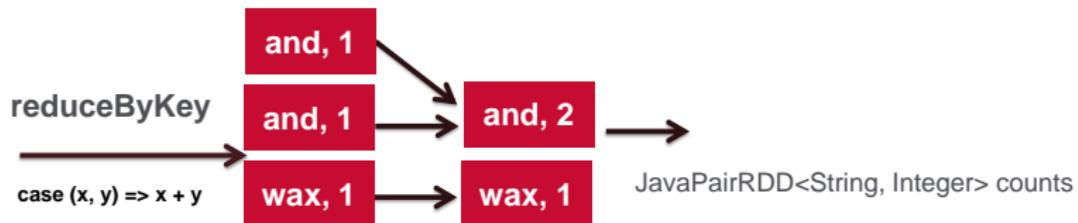
```
val input = sc.textFile(inputFile)                                HadoopRDD  
val words = input.flatMap(line => line.split(" "))              MapPartitionsRDD  
val counts = words  
    .map(word => (word, 1))                                         MapPartitionsRDD  
    .reduceByKey{case (x, y) => x + y}                                ShuffledRDD
```



Fonte: Carol McDonald: An Overview of Apache Spark



reduceByKey

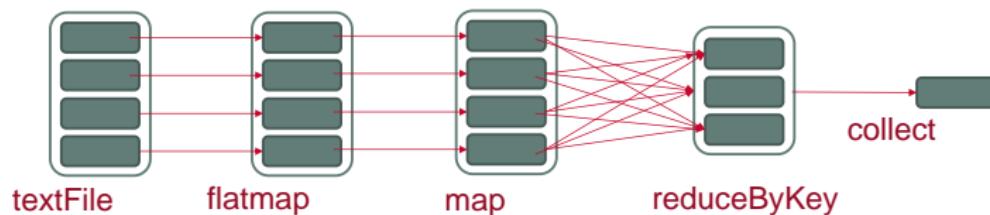


Fonte: Carol McDonald: An Overview of Apache Spark



Example Spark Word Count in Scala

```
val input = sc.textFile(inputFile)          HadoopRDD  
val words = input.flatMap(line => line.split(" ")) MapPartitionsRDD  
val counts = words  
    .map(word => (word, 1))  
    .reduceByKey{case (x, y) => x + y}  
val countArray = counts.collect()           ShuffledRDD  
                                            Array
```



Fonte: Carol McDonald: An Overview of Apache Spark



Transformations and Actions

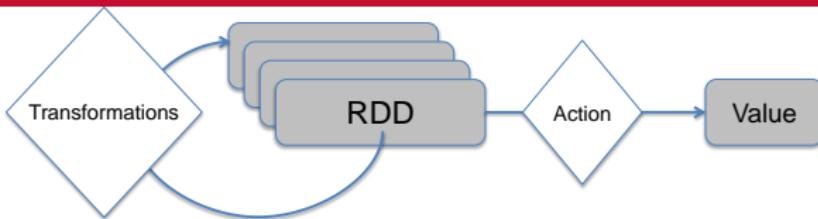


© 2014 MapR Technologies 

Fonte: Carol McDonald: An Overview of Apache Spark



RDD Transformations and Actions



Transformations
(define a new RDD)

- map
- filter
- sample
- union
- groupByKey
- reduceByKey
- join
- cache



Actions
(return a value)

- reduce
- collect
- count
- save
- lookupKey
- ...

Fonte: Carol McDonald: An Overview of Apache Spark



Basic Transformations

```
> nums = sc.parallelize([1, 2, 3])  
  
# Pass each element through a function  
> squares = nums.map(lambda x: x*x) // {1, 4, 9}  
  
# Keep elements passing a predicate  
> even = squares.filter(lambda x: x % 2 == 0) // {4}  
  
# Map each element to zero or more others  
> nums.flatMap(lambda x: range(x))  
  > # => {0, 0, 1, 0, 1, 2}
```

Range object (sequence of
numbers 0, 1, ..., x-1)



DATABRICKS

Fonte: Carol McDonald: An Overview of Apache Spark



Basic Actions

```
> nums = sc.parallelize([1, 2, 3])  
# Retrieve RDD contents as a local collection  
> nums.collect() # => [1, 2, 3]  
# Return first K elements  
> nums.take(2) # => [1, 2]  
# Count number of elements  
> nums.count() # => 3  
# Merge elements with an associative function  
> nums.reduce(lambda x, y: x + y) # => 6  
# Write elements to a text file  
> nums.saveAsTextFile("hdfs://file.txt")
```



Fonte: Carol McDonald: An Overview of Apache Spark



There's a lot more !



© 2014 MapR Technologies 

Fonte: Carol McDonald: An Overview of Apache Spark

Projeto 2

Neste projeto vamos comparar o uso do Hadoop MapReduce com o Apache Spark.

- ▶ Escolha uma aplicação
- ▶ Estude o código e proponha uma pequena alteração
- ▶ O modelo de transformações e ações foi mais adequado?



Referências

- ▶ <http://spark.apache.org/>
- ▶ An Overview of Apache Spark, Carol McDonald
- ▶ Clash of Titans: MapReduce vs. Spark for Large Scale Data Analytics, Juwei Shi e outros, IBM Research, China